

# EDA

July 21, 2025

## 1 COVID-19 County-Level Case Analysis

### Exploratory Data Analysis Report

Prepared by: Navid Bahadorn

Date: June, 2025

This report presents a preliminary analysis of COVID-19 county-level case data in the United States, including spatial and temporal trends.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
from config import config
import torch

sns.set(style="whitegrid")
pd.set_option('display.max_columns', 100)
```

#### 1.0.1 1. Dataset Summary

We first look at the dimensions and basic statistics of the dataset.

```
[2]: data_path = config["data_path"]
df = pd.read_csv(data_path)

# Restore dates from "t"
start_date = datetime(2020, 1, 21)
df["date"] = pd.to_datetime(df["t"], unit="D", origin=start_date)

print("Loaded shape:", df.shape)
display(df.head())
print(df.info())
```

Loaded shape: (2409352, 5)

	lon	lat	t	u	date
0	-164.040108	64.942121	85	10.167768	2020-04-15
1	-164.040108	64.942121	86	10.167768	2020-04-16

```

2 -164.040108  64.942121  87  10.167768 2020-04-17
3 -164.040108  64.942121  88  10.167768 2020-04-18
4 -164.040108  64.942121  89  10.167768 2020-04-19

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2409352 entries, 0 to 2409351
Data columns (total 5 columns):
#   Column  Dtype
---  -
0    lon    float64
1    lat    float64
2     t     int64
3     u     float64
4   date    datetime64[ns]
dtypes: datetime64[ns](1), float64(3), int64(1)
memory usage: 91.9 MB
None

```

## 1.0.2 2. Summary Statistics & Nulls

```

[3]: print("Summary Statistics:")
      display(df.describe())

      print("Null Value Check:")
      print(df.isnull().sum())

```

Summary Statistics:

	lon	lat	t	u \
count	2.409352e+06	2.409352e+06	2.409352e+06	2.409352e+06
mean	-9.214093e+01	3.838049e+01	4.573247e+02	1.011351e+04
min	-1.640401e+02	1.960199e+01	0.000000e+00	1.028686e-02
25%	-9.808979e+01	3.464985e+01	2.650000e+02	2.138229e+03
50%	-9.024324e+01	3.832408e+01	4.580000e+02	9.257088e+03
75%	-8.343685e+01	4.174292e+01	6.510000e+02	1.522107e+04
max	-6.762834e+01	6.935328e+01	8.430000e+02	3.843137e+05
std	1.261735e+01	5.230411e+00	2.235216e+02	8.760030e+03

	date
count	2409352
mean	2021-04-22 07:47:36.891645440
min	2020-01-21 00:00:00
25%	2020-10-12 00:00:00
50%	2021-04-23 00:00:00
75%	2021-11-02 00:00:00
max	2022-05-13 00:00:00
std	NaN

Null Value Check:

```
lon      0
lat      0
t        0
u        0
date     0
dtype: int64
```

### 1.0.3 3.Unique Locations & Data Coverage

```
[5]: print(f"Unique locations (lon, lat): {df[['lon','lat']].drop_duplicates().
      ↪shape[0]}")
      print(f"Time Range: {df['date'].min().date()} to {df['date'].max().date()}")
```

```
Unique locations (lon, lat): 3124
Time Range: 2020-01-21 to 2022-05-13
```

#### Observation:

The dataset includes infection rate per 100,000 people (u) over time for each U.S. county. The data ranges from Jan 2020 onwards and appears to be normalized by population.


---

### 1.0.4 4.Sample Time Series for a Random or Top County

We examine the infection trend over time for a sample location with the most data points.

```
[6]: loc = df.groupby(["lon", "lat"]).size().sort_values(ascending=False).index[0]
      df_sample = df[(df["lon"] == loc[0]) & (df["lat"] == loc[1])]

      plt.figure(figsize=(10, 4))
      plt.plot(df_sample["date"], df_sample["u"], marker="o")
      plt.title("Infection Rate Over Time at a Sample County")
      plt.xlabel("Date")
      plt.ylabel("Cases per 100k")
      plt.grid(True)
      plt.tight_layout()
      plt.show()
```



### Analysis:


This county shows distinct surges during known COVID waves (e.g., early 2021, late 2021/Omicron). The infection rate exhibits strong temporal spikes followed by rapid declines, consistent with public health intervention cycles.

#### 1.0.5 5. National Average Trend

Here, we analyze the average infection rate across all counties each day.

```
[7]: df_national = df.groupby("date")["u"].mean().reset_index()

plt.figure(figsize=(10, 4))
sns.lineplot(data=df_national, x="date", y="u", color="coral")
plt.title("National Average Cases per 100k Over Time")
plt.xlabel("Date")
plt.ylabel("Avg Infection Rate")
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```



### Analysis:


National averages smooth out county-level noise and highlight broad pandemic waves. The three major surges are clearly visible, corresponding to initial outbreak, Delta, and Omicron waves.

## 1.0.6 6. Heatmap: Top Locations

We now plot trends for the top 5 counties by data density.

```
[8]: # Get top 5 locations with most records
top5_locs = df.groupby(["lon", "lat"]).size().sort_values(ascending=False).
    ↪head(5).index

# Plot time series for each
plt.figure(figsize=(12, 6))
for i, loc in enumerate(top5_locs):
    df_sub = df[(df["lon"] == loc[0]) & (df["lat"] == loc[1])]
    plt.plot(df_sub["date"], df_sub["u"], label=f"Loc {i+1}")
plt.legend()
plt.title("Infection Rate Trends - Top 5 Locations")
plt.xlabel("Date")
plt.ylabel("Cases per 100k")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



### Analysis:


Comparing different counties reveals variations in timing and magnitude of infection waves. Some counties show delayed or prolonged peaks, possibly due to differences in population density, interventions, or mobility.

#### 1.0.7 Heatmap of a specific time slice

```
[7]: df_t = df[df["t"] == df["t"].max()].copy()
df_t["lat_round"] = df_t["lat"].round(1)
df_t["lon_round"] = df_t["lon"].round(1)

heatmap_data = df_t.pivot_table(index="lat_round", columns="lon_round",
    ↪ values="u", aggfunc="mean")
heatmap_data = heatmap_data.fillna(0)

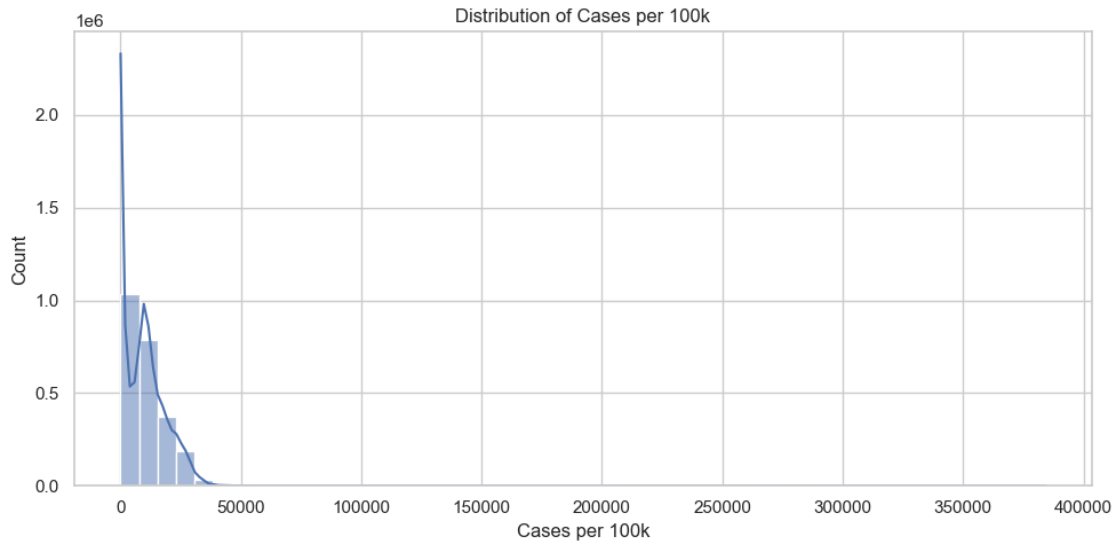
# Optionally log-scale for visualization
plt.figure(figsize=(10, 8))
sns.heatmap(np.log1p(heatmap_data), cmap="Reds")
plt.title("Spatial Distribution of log(Cases + 1)")
plt.xlabel("Longitude")
plt.ylabel("Latitude")
plt.tight_layout()
plt.show()
```



### 1.0.8 7.Daily Histogram of Cases per 100k

Next, we examine how infection rates are distributed overall.

```
[23]: plt.figure(figsize=(10, 5))
sns.histplot(df["u"], bins=50, kde=True)
plt.title("Distribution of Cases per 100k")
plt.xlabel("Cases per 100k")
plt.tight_layout()
plt.show()
```



### Analysis:

The distribution is right-skewed, with a long tail of counties experiencing very high infection rates. Most data points fall below 500 per 100k.

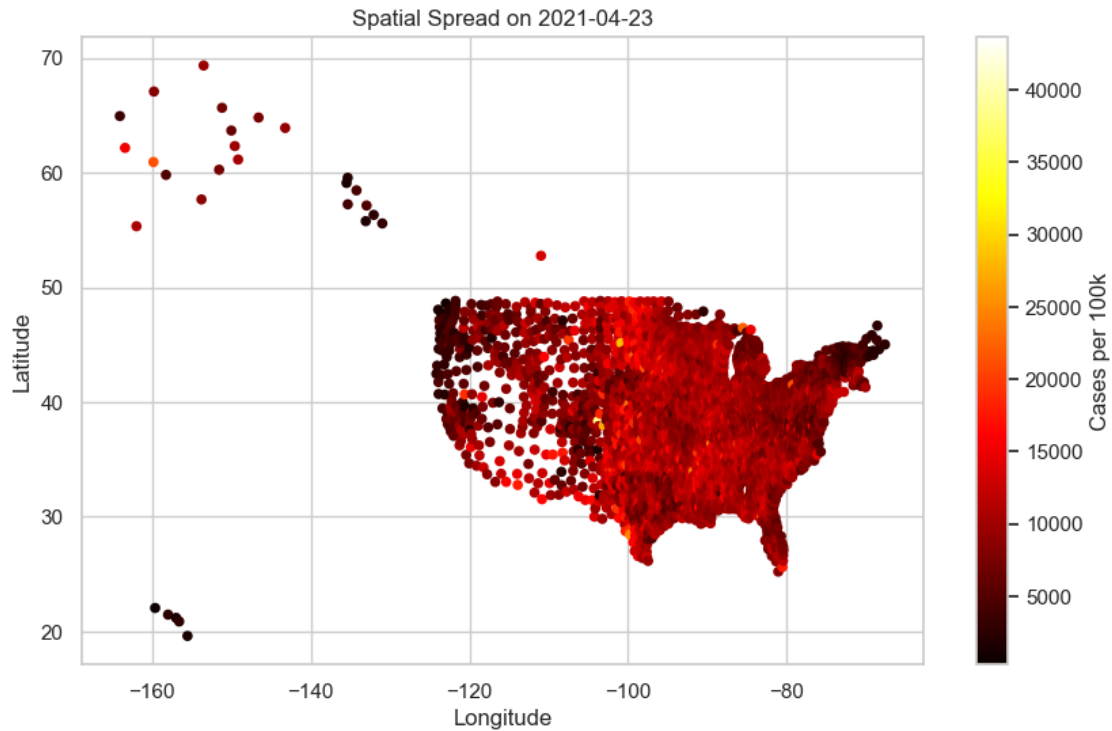
## 1.0.9 8.Geographic Snapshots for a Specific Date

We visualize the infection rate across counties on a single median day.

```
[25]: # Pick a snapshot day (e.g., peak period)
snapshot_day = df["date"].median()
df_day = df[df["date"] == snapshot_day]

plt.figure(figsize=(10, 6))
plt.scatter(df_day["lon"], df_day["lat"], c=df_day["u"], cmap="hot", s=20)
plt.colorbar(label="Cases per 100k")
plt.title(f"Spatial Spread on {snapshot_day.date()}")
plt.xlabel("Longitude")
plt.ylabel("Latitude")
plt.grid(True)
plt.show()
```





### Analysis:


Geographic variation in infection intensity is clearly visible. Certain hotspots are evident, possibly urban centers. This spatial map offers a useful way to model regional dynamics later.

#### 1.0.10 9.Rolling Average

This smooths daily noise and captures trends.

```
[26]: df_rolling = df.groupby("date")["u"].mean().rolling(window=14).mean().
      ↪reset_index()

plt.figure(figsize=(10, 4))
sns.lineplot(data=df_rolling, x="date", y="u", color="blue")
plt.title("14-Day Rolling Avg Infection Rate (National)")
plt.xlabel("Date")
plt.ylabel("Smoothed Cases per 100k")
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```




### Analysis:

This plot smooths out day-to-day volatility, showing the three major pandemic phases with clear transitions and drop-offs. This can help guide training windows for time-series models.

```
[27]: # Pick a random location
sample = df.groupby(["lon", "lat"]).size().sort_values(ascending=False).index[0]
df_sample = df[(df["lon"] == sample[0]) & (df["lat"] == sample[1])]

plt.figure(figsize=(10, 4))
plt.plot(df_sample["t"], df_sample["u"], label="Normalized Cases")
plt.title(f"u(t) at County (lon={sample[0]:.2f}, lat={sample[1]:.2f})")
plt.xlabel("Time t (days)"); plt.ylabel("u (per 100k)")
plt.grid(True)
plt.legend()
plt.show()
```



### 1.0.11 Conclusions

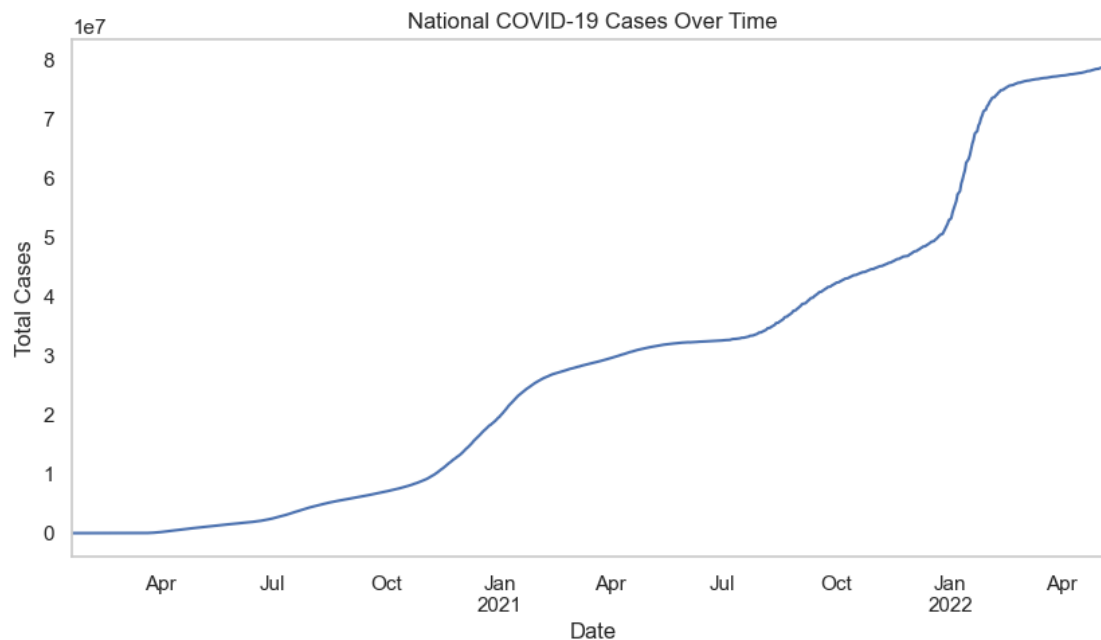
- Temporal trends align with known pandemic waves.
- Spatial variation suggests that geographic modeling is necessary.
- Rolling averages are essential to suppress noise in training data.

This analysis supports using Physics-Informed Neural Networks (PINNs) with both spatial and temporal inputs for modeling spread dynamics.

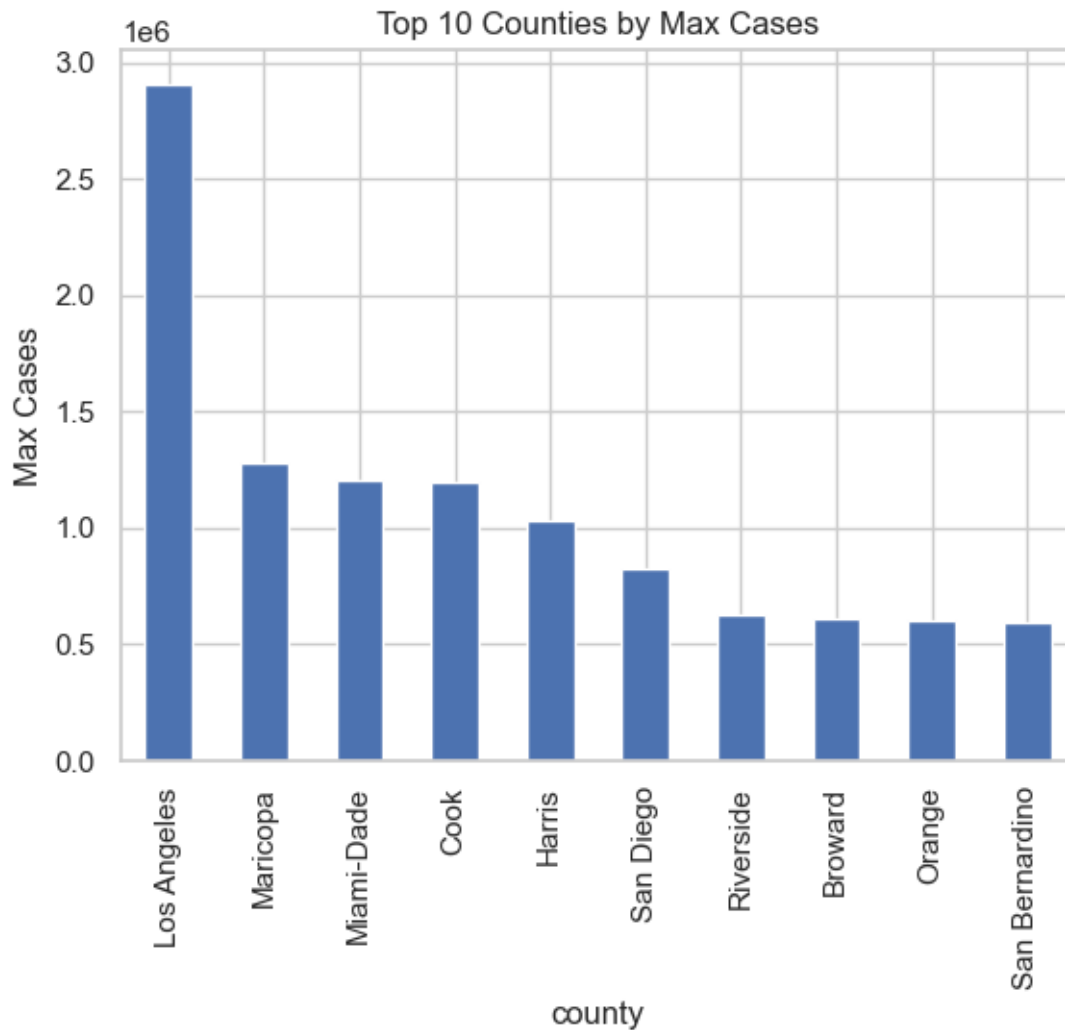
---

```
[29]: # Convert date
df["date"] = pd.to_datetime(df["date"])
```

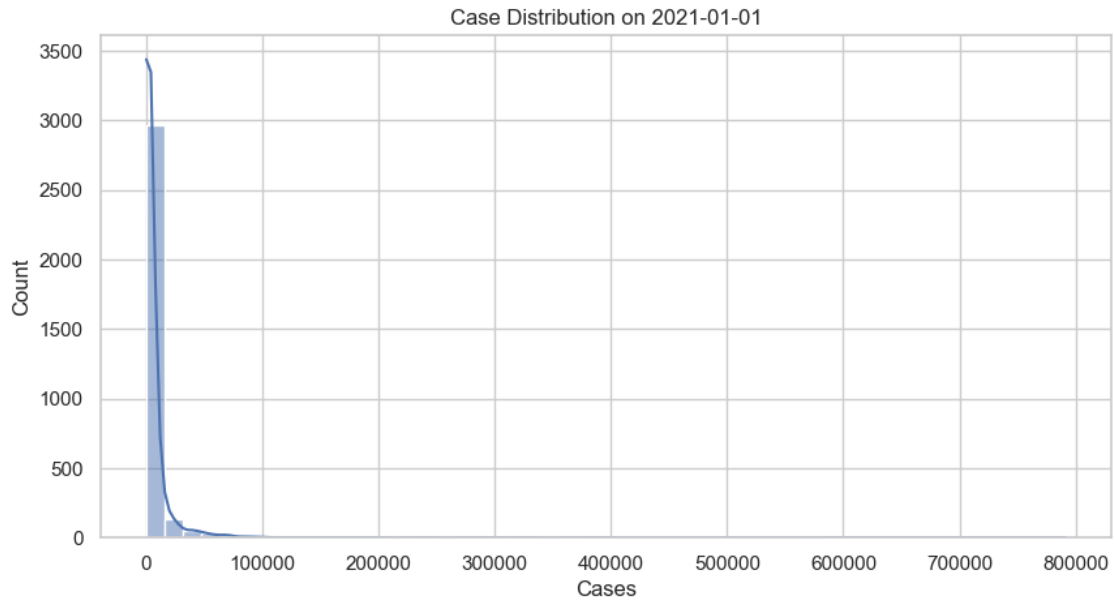
```
[30]: # Total cases over time (U.S.)
national_cases = df.groupby("date")["cases"].sum()
plt.figure(figsize=(10, 5))
national_cases.plot(title="National COVID-19 Cases Over Time")
plt.xlabel("Date")
plt.ylabel("Total Cases")
plt.grid()
plt.show()
```



```
[31]: # Top 10 counties by total cases
top_counties = df.groupby("county")["cases"].max().nlargest(10)
top_counties.plot(kind="bar", title="Top 10 Counties by Max Cases")
plt.ylabel("Max Cases")
plt.show()
```



```
[32]: # Cases heatmap: example on a single date
single_day = df[df["date"] == "2021-01-01"]
pivot = single_day.pivot_table(index="fips", values="cases")
plt.figure(figsize=(10, 5))
sns.histplot(single_day["cases"], bins=50, kde=True)
plt.title("Case Distribution on 2021-01-01")
plt.xlabel("Cases")
plt.show()
```



```
[33]: data_path = config["data_path"]
pinn_df = pd.read_csv(data_path)

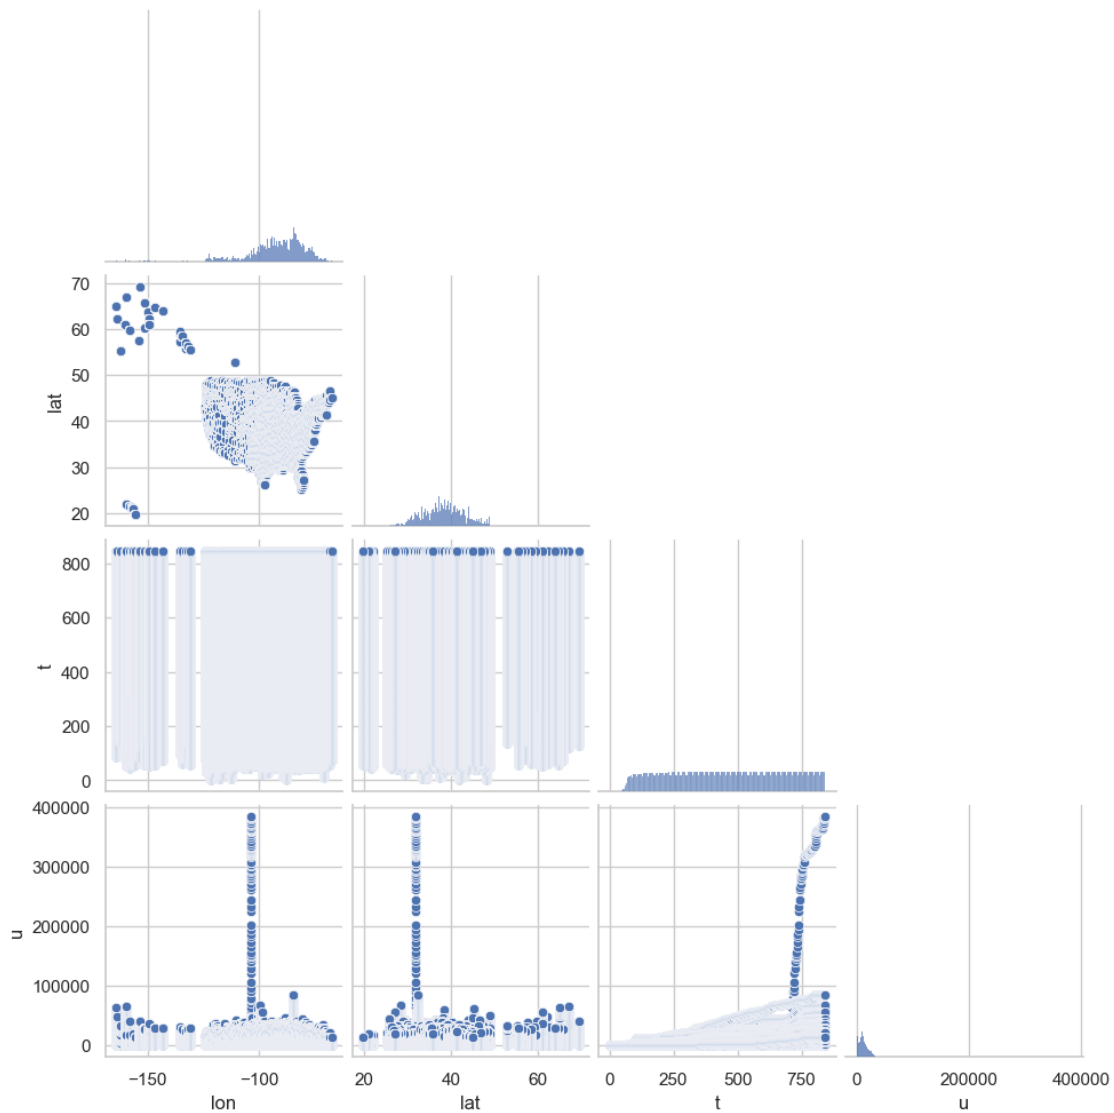
print("[INFO] PINN Dataset Summary")
display(pinn_df.describe())

# Pairplot of inputs
sns.pairplot(pinn_df[["lon", "lat", "t", "u"]], corner=True)
plt.suptitle("Pairplot of PINN Inputs (log-scaled u)", y=1.02)
plt.show()
```

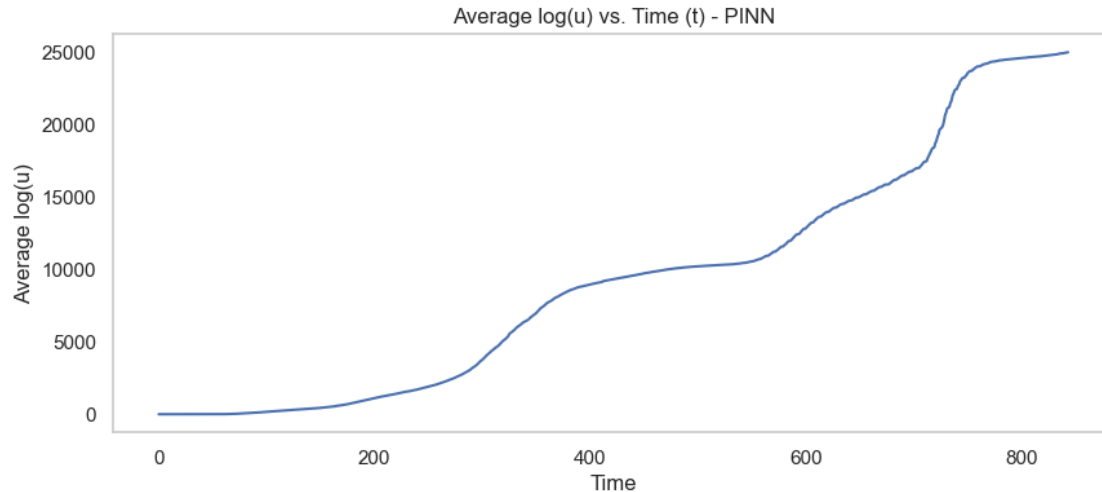
[INFO] PINN Dataset Summary

	lon	lat	t	u
count	2.409352e+06	2.409352e+06	2.409352e+06	2.409352e+06
mean	-9.214093e+01	3.838049e+01	4.573247e+02	1.011351e+04
std	1.261735e+01	5.230411e+00	2.235216e+02	8.760030e+03
min	-1.640401e+02	1.960199e+01	0.000000e+00	1.028686e-02
25%	-9.808979e+01	3.464985e+01	2.650000e+02	2.138229e+03
50%	-9.024324e+01	3.832408e+01	4.580000e+02	9.257088e+03
75%	-8.343685e+01	4.174292e+01	6.510000e+02	1.522107e+04
max	-6.762834e+01	6.935328e+01	8.430000e+02	3.843137e+05

Pairplot of PINN Inputs (log-scaled u)



```
[36]: # Time series average
avg_u_time = pinn_df.groupby("t")["u"].mean()
plt.figure(figsize=(10, 4))
avg_u_time.plot()
plt.title("Average log(u) vs. Time (t) - PINN")
plt.xlabel("Time")
plt.ylabel("Average log(u)")
plt.grid()
plt.show()
```



```
[40]: # Unpack
graph_dataset, node_positions = torch.load(config["gnn_dataset_path"],
      ↪weights_only=False)

print(f"[INFO] Loaded {len(graph_dataset)} graph snapshots.")

# Access first graph snapshot
first_graph = graph_dataset[0]

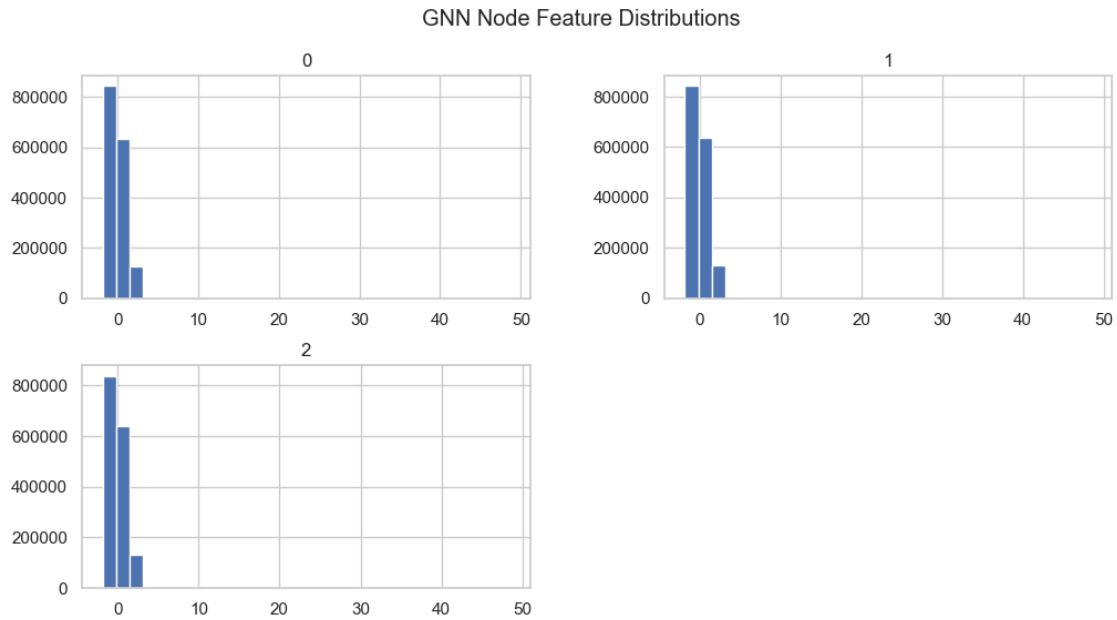
# EDA
num_nodes = first_graph.num_nodes
num_edges = first_graph.edge_index.shape[1]
avg_deg = num_edges / num_nodes

print(f" GNN Graph Summary:")
print(f"   • Nodes: {num_nodes}")
print(f"   • Edges: {num_edges}")
print(f"   • Avg Degree: {avg_deg:.2f}")
```

```
[INFO] Loaded 516 graph snapshots.
GNN Graph Summary:
  • Nodes: 3124
  • Edges: 10410
  • Avg Degree: 3.33
```


```
[41]: # Feature & target stats
x_all = torch.cat([g.x for g in graph_dataset], dim=0).numpy()
y_all = torch.cat([g.y for g in graph_dataset], dim=0).numpy()
```

```
[42]: # Node feature histogram
pd.DataFrame(x_all).hist(bins=30, figsize=(12, 6))
plt.suptitle("GNN Node Feature Distributions")
plt.show()
```



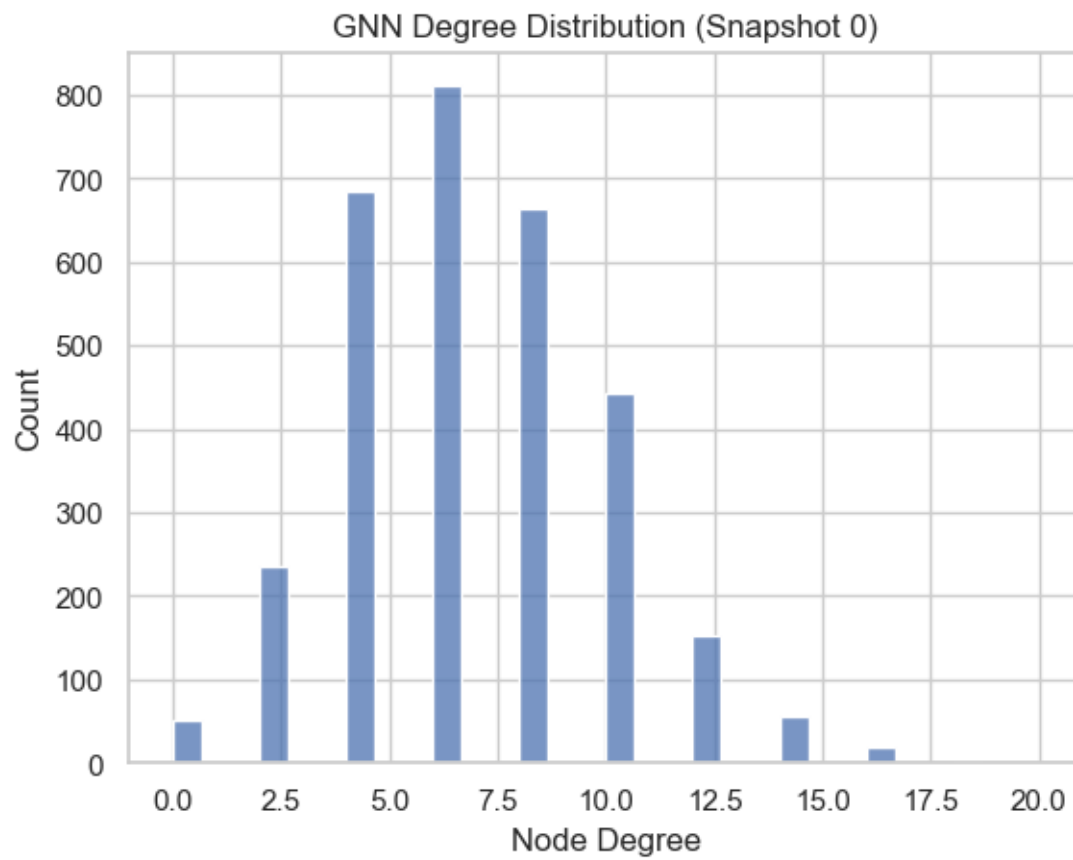
```
[43]: # Target histogram
sns.histplot(y_all, bins=50, kde=True)
plt.title("GNN Target (y) Distribution Across Snapshots")
plt.xlabel("log(cases + 1)")
plt.show()
```





```
[44]: # Degree distribution of a snapshot
from torch_geometric.utils import to_networkx
import networkx as nx

G = to_networkx(graph_dataset[0])
degrees = [deg for _, deg in G.degree()]
sns.histplot(degrees, bins=30)
plt.title("GNN Degree Distribution (Snapshot 0)")
plt.xlabel("Node Degree")
plt.ylabel("Count")
plt.show()
```



```
[45]: # Time slice trends (target mean over snapshots)
y_per_snapshot = [g.y.mean().item() for g in graph_dataset]
plt.plot(range(len(graph_dataset)), y_per_snapshot)
plt.title("Mean log(y) per Graph Snapshot (GNN)")
plt.xlabel("Time Index")
plt.ylabel("Mean log(y)")
plt.grid()
plt.show()
```

